

# Lab 5: Inter-Process Communication (Part II)

Yu-Lun Huang

2011-11-21

## 1 Objective

- Be familiar with inter-process communication: pipe, shared memory, etc.

## 2 Prerequisite

- Read man pages of `pipe()`, `shmget()`, `shmopt()`, `shmat()`, `shmdt()`, etc.

## 3 Pipe

A pipe creates a pair of file descriptors, pointing to a pipe inode, and places them in the array pointed to by `filedes`. `filedes[0]` is for reading, `filedes[1]` is for writing. An unnamed pipe can be viewed and accessed by processes with parent-child relationship. To share a pipe among all processes, you need to create a named pipe.

The following program creates a pipe, and then `fork(2)` to create a child process. After the `fork(2)`, each process closes the descriptors that it doesn't need for the pipe (see `pipe(7)`). In the following sample code, the child process reads the file specified in `argv[1]`, and writes the file content to the parent process through the pipe. The parent process reads the data from the pipe and echoes the data on the screen.

```
1  /* pipe.c
2  *
3  * child process read the content of file
4  * and write the content to parent process through pipe
5  */
6
7  #include <fcntl.h>
8  #include <stdio.h>
9  #include <stdlib.h>
10 #include <string.h>
11 #include <sys/stat.h>
12 #include <sys/types.h>
13 #include <sys/wait.h>
14 #include <unistd.h>
15
16 int pfd[2]; /* pfd[0] is read end, pfd[1] is write end */
17
18 void ChildProcess(char *path)
19 {
20     int fd;
```

```

21  int ret;
22  char buffer[100];
23
24  /* close unused read end */
25  close(pfd[0]);
26
27  /* open file */
28  fd = open(path, ORDONLY);
29  if (fd < 0) {
30      printf("Open_%s_failed.\n", path);
31      exit(EXIT_FAILURE);
32  }
33
34  /* read file and write content to pipe */
35  while (1) {
36      /* read raw data from file */
37      ret = read(fd, buffer, 100);
38
39      if (ret < 0) { /* error */
40          perror("read()");
41          exit(EXIT_FAILURE);
42      }
43      else if (ret == 0) { /* reach EOF */
44          close(fd); /* close file */
45          close(pfd[1]); /* close write end, reader see EOF */
46          exit(EXIT_SUCCESS);
47      }
48      else { /* write content to pipe */
49          write(pfd[1], buffer, ret);
50      }
51  }
52 }
53
54 void ParentProcess()
55 {
56     int ret;
57     char buffer[100];
58
59     /* close unused write end */
60     close(pfd[1]);
61
62     /* read data from pipe until reach EOF */
63     while(1) {
64         ret = read(pfd[0], buffer, 100);
65
66         if (ret > 0) { /* print data to screen */
67             printf("%.s", ret, buffer);
68         }
69         else if (ret == 0) { /* reach EOF */
70             close(pfd[0]); /* close read end */
71             wait(NULL);

```

```

72         exit(EXIT_SUCCESS);
73     }
74     else {
75         perror("pipe_read()");
76         exit(EXIT_FAILURE);
77     }
78 }
79 }
80
81 int main(int argc, char *argv[])
82 {
83     pid_t cpid;
84
85     if (argc != 2) {
86         fprintf(stderr, "%s: specify a file\n", argv[0]);
87         exit(1);
88     }
89
90     /* create pipe */
91     if (pipe(pfd) == -1) {
92         perror("pipe");
93         exit(EXIT_FAILURE);
94     }
95
96     /* fork child process */
97     cpid = fork();
98     if (cpid == -1) { /* error */
99         perror("fork");
100         exit(EXIT_FAILURE);
101     }
102
103     if (cpid == 0)
104         ChildProcess(argv[1]);
105     else
106         ParentProcess();
107
108     return 0;
109 }

```

Monitoring multiple file descriptors with polling strategy may cause busy waiting, which lowers down the system performance. To effectively monitor multiple descriptors, you can use the system call `select()` to perform block waiting, rather than busy waiting. The calling process specifies the interesting file descriptors to `select()` and performs blocking waiting. When one or more descriptors become "ready" (ready for read or ready for write), `select()` informs the calling process to awake from blocking state. Then, a user can check each file descriptor to perform read/write operation. In following program, two child processes sleep a random time, then send a message to the parent process. The parent process uses `select()` to perform blocking waiting, until one of the pipes is ready to read.

```

1  /*
2   * select.c
3   */
4

```

```

5 #include <stdio.h>
6 #include <stdlib.h>
7 #include <string.h>
8 #include <sys/time.h>
9 #include <sys/types.h>
10 #include <time.h>
11 #include <unistd.h>
12
13 #define max(a, b) ((a > b) ? a : b)
14
15 void ChildProcess(int *pfd, int sec)
16 {
17     char buffer[100];
18
19     /* close unused read end */
20     close(pfd[0]);
21
22     /* sleep a random time to wait parent process enter select() */
23     printf("Child_process_(%d)_wait_%d_secs\n", getpid(), sec);
24     sleep(sec);
25
26     /* write message to parent process */
27     memset(buffer, 0, 100);
28     sprintf(buffer, "Child_process_(%d)_sent_message_to_parent_process\n",
29             getpid());
30     write(pfd[1], buffer, strlen(buffer));
31
32     /* close write end */
33     close(pfd[1]);
34
35     exit(EXIT_SUCCESS);
36 }
37
38 int main(int argc, char *argv[])
39 {
40     int pfd1[2], pfd2[2]; /* pipe's fd */
41     int cpid1, cpid2; /* child process id */
42     fd_set rfd, arfd;
43     int max_fd;
44     struct timeval tv;
45     int retval;
46     int fd_index;
47     char buffer[100];
48
49     /* random seed */
50     srand(time(NULL));
51
52     /* create pipe */
53     pipe(pfd1);
54     pipe(pfd2);
55

```

```

56  /* create 2 child processes and set corresponding pipe & sleep time */
57  cpid1 = fork();
58  if (cpid1 == 0)
59      ChildProcess(pfd1, random() % 5);
60
61  cpid2 = fork();
62  if (cpid2 == 0)
63      ChildProcess(pfd2, random() % 4);
64
65  /* close unused write end */
66  close(pfd1[1]);
67  close(pfd2[1]);
68
69  /* set pfd1[0] & pfd2[0] to watch list */
70  FD_ZERO(&rfd);
71  FD_ZERO(&rfd);
72  FD_SET(pfd1[0], &rfd);
73  FD_SET(pfd2[0], &rfd);
74  max_fd = max(pfd1[0], pfd2[0]) + 1;
75
76  /* Wait up to five seconds. */
77  tv.tv_sec = 5;
78  tv.tv_usec = 0;
79
80  while(1)
81  {
82      /* config fd_set for select */
83      memcpy(&rfd, &rfd, sizeof(rfd));
84
85      /* wait until any fd response */
86      retval = select(max_fd, &rfd, NULL, NULL, &tv);
87
88      if (retval == -1) { /* error */
89          perror("select()");
90          exit(EXIT_FAILURE);
91      }
92      else if (retval) { /* # of fd got response */
93          printf("Data is available now.\n");
94      }
95      else { /* no fd response before timer expired */
96          printf("No data within five seconds.\n");
97          break;
98      }
99
100     /* check if any response */
101     for (fd_index = 0; fd_index < max_fd; fd_index++)
102     {
103         if (!FD_ISSET(fd_index, &rfd))
104             continue; /* no response */
105
106         retval = read(fd_index, buffer, 100);

```

```

107
108         if (retval > 0)          /* read data from pipe */
109             printf("%.s", retval, buffer);
110         else if (retval < 0) /* error */
111             perror("pipe_read()");
112         else {                  /* write fd closed */
113             /* close read fd */
114             close(fd_index);
115             /* remove fd from watch list */
116             FD_CLR(fd_index, &arfds);
117         }
118     }
119 }
120
121 return 0;
122 }

```

## 4 Shared Memory

Shared memory is a memory space that may be simultaneously accessed by multiple processes with an intent to provide inter-process communication among them or avoid redundant copies. Unlike unnamed pipes, only exist among processes with parent-child relationship, every process can access the shared memory space with the **share memory key** specified.

The followings are two processes communicating via shared memory: **shm\_server.c** and **shm\_client.c**. The two programs here illustrate the passing of a simple piece of memory (a string) between the processes if running simultaneously:

```

1  /*
2  * shm_server.c -- creates the string and shared memory.
3  */
4
5  #include <stdio.h>
6  #include <stdlib.h>
7  #include <sys/ipc.h>
8  #include <sys/shm.h>
9  #include <sys/types.h>
10 #include <unistd.h>
11
12 #define SHMSZ      27
13
14 int main(int argc, char *argv[])
15 {
16     char c;
17     int shmid;
18     key_t key;
19     char *shm, *s;
20     int retval;
21
22     /* We'll name our shared memory segment "5678" */
23     key = 5678;
24

```

```

25  /* Create the segment */
26  if ((shm = shmget(key, SHMSZ, IPC_CREAT | 0666)) < 0) {
27      perror("shmget");
28      exit(1);
29  }
30
31  /* Now we attach the segment to our data space */
32  if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
33      perror("shmat");
34      exit(1);
35  }
36  printf("Server_create_and_attach_the_share_memory.\n");
37
38  /* Now put some things into the memory for the other process to read */
39  s = shm;
40
41  printf("Server_write_a_~_z_to_share_memory.\n");
42  for (c = 'a'; c <= 'z'; c++)
43      *s++ = c;
44  *s = '\0';
45
46  /*
47   * Finally, we wait until the other process changes the first
48   * character of our memory to '*', indicating that it has read
49   * what we put there.
50   */
51  printf("Waiting_other_process_read_the_share_memory...\n");
52  while (*shm != '*')
53      sleep(1);
54  printf("Server_read_*_from_the_share_memory.\n");
55
56  /* Detach the share memory segment */
57  shmdt(shm);
58
59  /* Destroy the share memory segment */
60  printf("Server_destroy_the_share_memory.\n");
61  retval = shmctl(shmid, IPC_RMID, NULL);
62  if (retval < 0)
63  {
64      fprintf(stderr, "Server_remove_share_memory_failed\n");
65      exit(1);
66  }
67
68  return 0;
69 }

```

```

1  /*
2   * shm_client.c -- attaches itself to the created shared memory
3   *                  and uses the string (printf).
4   */
5

```

```

6 #include <stdio.h>
7 #include <stdlib.h>
8 #include <sys/ipc.h>
9 #include <sys/shm.h>
10 #include <sys/types.h>
11
12 #define SHMSZ      27
13
14 int main(int argc, char *argv[])
15 {
16     int shmid;
17     key_t key;
18     char *shm, *s;
19
20     /* We need to get the segment named "5678", created by the server */
21     key = 5678;
22
23     /* Locate the segment */
24     if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
25         perror("shmget");
26         exit(1);
27     }
28
29     /* Now we attach the segment to our data space */
30     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
31         perror("shmat");
32         exit(1);
33     }
34     printf("Client_attach_the_share_memory_created_by_server.\n");
35
36     /* Now read what the server put in the memory */
37     printf("Client_read_characters_from_share_memory...\n");
38     for (s = shm; *s != '\0'; s++)
39         putchar(*s);
40     putchar('\n');
41
42     /*
43      * Finally, change the first character of the segment to '*',
44      * indicating we have read the segment.
45      */
46     printf("Client_write_*_to_the_share_memory.\n");
47     *shm = '*';
48
49     /* Detach the share memory segment */
50     printf("Client_detach_the_share_memory.\n");
51     shmdt(shm);
52
53     return 0;
54 }

```



# Lab 5 Assignments

Yu-Lun Huang

2011-11-21

- Rewrite your Lab3 assignment and use **pipe** to transfer LED control signals between processes. (Note: Do not create a child process for each control signal)
- Rewrite your Lab4 assignment and replace the two files with **shared memories**.
- Design a program to calculate the summation of a series.
  - The program allows user to input two integers, **N** and **M**.
  - The parent process then creates **N** child processes to calculate the summation from **1** to **N×M**
  - Each child process receives a serial number **n** (ranged from 0 to  $N - 1$ ) and the integer **M** from parent process.
  - The child process calculates  $\sum_{(nM+1)}^{(nM+M)}$  and writes the summation result back to parent process using **pipe**.
  - The parent process uses **select()** to read the summation results of child processes and sums up the final result.
  - Display the last 4 digits of the final result on 7-segment display.